

```
In [1]: # Created by: Jessica Gallo
# Date Created: 5/7/2021
# Last Modified: 5/13/2021
# SVM for Face Recognition/Classification
# RGB -> Gaussian Filter -> Histogram Equalization -> Sobel Filter
```

```
In [2]: # -----
# GPU |
# ----

import tensorflow as tf
print(tf.test.gpu_device_name())
import tensorflow
print(tensorflow.__version__)
import keras
print(keras.__version__)
```

2.1.0

2.3.1

Using TensorFlow backend.

```
In [3]: # -----
# IMPORTS |
# -----
import os, shutil
import glob
import pandas as pd
import matplotlib.pyplot as plt
from tensorflow.keras import layers
from tensorflow.keras import models
from tensorflow.keras import optimizers
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras import optimizers
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from keras.preprocessing import image # preprocessing a single image
import numpy as np # preprocessing a single image
from keras.applications import VGG16 # defining the loss tensor for filter visualization
from keras import backend as K # defining the loss tensor for filter visualization
import random
from sklearn import metrics
from sklearn.metrics import confusion_matrix
from sklearn.metrics import roc_curve, auc
from sklearn.metrics import roc_auc_score
from sklearn.metrics import classification_report
import seaborn as sns
```

```
In [4]: # -----  
# LOAD Dataset |  
# -----  
  
# 113 total pictures in each folder  
# 91 total pictures in each folder for training  
# 22 total pictures in each folder for testing  
# Path to dataset  
  
dataset_dir = './Documents/MATLAB/RGB_Enhanced_Dataset'  
  
os.listdir('./Documents/MATLAB/RGB_Enhanced_Dataset') # returns list
```

```
Out[4]: ['test', 'train']
```

```
In [5]: # -----  
# Mapping directories |  
# -----  
  
train_dir = './Documents/MATLAB/RGB_Enhanced_Dataset/train'  
test_dir = './Documents/MATLAB/RGB_Enhanced_Dataset/test'
```

```
In [6]: # -----  
# Listing Amount of Images for Each Expression |  
# -----  
  
train_count=[]  
val_count=[]  
whole_count=[]  
  
print('TRAINING SET')  
files= os.listdir("../Documents//MATLAB//RGB_Enhanced_Dataset//train")  
for type in files:  
    count = os.listdir("../Documents//MATLAB//RGB_Enhanced_Dataset//train//'+type+'")  
    print(type+ " " + str(len(count)))  
    train_count.append(len(count))  
print()  
  
print('TEST SET')  
files= os.listdir("../Documents//MATLAB//RGB_Enhanced_Dataset//test")  
  
for type in files:  
    count = os.listdir("../Documents//MATLAB//RGB_Enhanced_Dataset//test//'+type+'")  
    print(type+ " " + str(len(count)))  
    whole_count.append(len(count))
```

```
TRAINING SET  
1_Neutral 89  
2_Smiling 88  
3_Sleepy 89  
4_Surprise 89  
5_Sunglasses 89
```

```
TEST SET  
1_Neutral 22  
2_Smiling 22  
3_Sleepy 22  
4_Surprise 22  
5_Sunglasses 22
```

```
In [7]: # -----  
# Number of Samples in Each Directory |  
# -----  
  
# returns a number of items inside the folder  
def getNumber(path):  
    s = 0  
    for i in os.listdir(path):  
        if i != '.DS_Store':  
            s += len(os.listdir(os.path.join(path,i)))  
    return s  
  
n_train = getNumber(train_dir)  
n_test = getNumber(test_dir)  
  
print('Number of Samples Train:', n_train)  
print('Number of Samples Test:', n_test)
```



Number of Samples Train: 444
Number of Samples Test: 110


```
In [8]: # -----
# Plots to Show Data Distribution/Amounts |
# -----

# General Datasets
fig, ax = plt.subplots()
ax.bar(['Train', 'Test'], [n_train, n_test])
plt.show()

# Each expression aracter in training set
plt.barh(files, train_count,)
plt.title('Training Dataset')
plt.show()

# Each expression in test set
plt.barh(files, whole_count,)
plt.title('Test Set')
plt.show()
```





```
In [9]: # -----
# ImageDataGenerator /
# -----

# TEST
test_datagen = ImageDataGenerator(rescale=1./255)

# TRAIN
train_datagen = ImageDataGenerator(rescale= 1./255,
                                   rotation_range = 40,
                                   width_shift_range = 0.2,
                                   height_shift_range = 0.2,
                                   shear_range = 0.2,
                                   horizontal_flip = True)
train_generator = train_datagen.flow_from_directory(train_dir,
                                                    target_size = (128, 128),
                                                    batch_size = 32,
                                                    class_mode = 'categorical')

# VALIDATION
validation_datagen = ImageDataGenerator(rescale= 1./255,
                                       rotation_range = 40,
                                       width_shift_range = 0.2,
                                       height_shift_range = 0.2,
                                       shear_range = 0.2,
                                       horizontal_flip = True)
validation_generator = test_datagen.flow_from_directory(test_dir,
                                                        target_size = (128, 128),
                                                        batch_size = 32,
                                                        class_mode = 'categorical',
                                                        shuffle=False)
```

Found 444 images belonging to 5 classes.
Found 110 images belonging to 5 classes.

```

In [10]: # -----
# Displaing Randomly Augmented Images |
# -----

from keras.preprocessing import image
train_dir = '../Documents//MATLAB//RGB_Enhanced_Dataset//train//1_Neutral'

fnames = [os.path.join(train_dir, fname) for
           fname in os.listdir(train_dir)]

img_path = fnames[6] #choosing an image to augment

img = image.load_img(img_path, target_size = (128, 128))

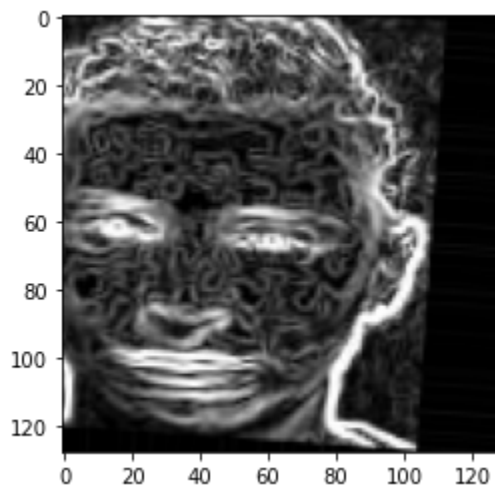
x = image.img_to_array(img)




x = x.reshape((1,) + x.shape)

i = 0
for batch in train_datagen.flow(x, batch_size=1):
    plt.figure(i)
    imgplot = plt.imshow(image.array_to_img(batch[0]))
    i+=1
    if i%4 == 0:
        break

plt.show()

```






```
In [14]: # MODEL 1
# -----
# ReLU activation function
# Adam optimizer
# 4 Conv
# 1 Batch Norm
# 4 MaxPool
# 1 Dropout
# 3 Dense
# 1 Flatten
# Batchsize 32

model = tensorflow.keras.Sequential()
model.add(layers.Conv2D(32, (3,3), activation='relu',
                        input_shape = (128, 128, 3),
                        kernel_initializer = 'glorot_normal',
                        bias_initializer = 'zeros'))
model.add(layers.BatchNormalization())
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(64, (3,3), activation='relu'))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(128, (3,3), activation='relu'))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Conv2D(128, (3,3), activation='relu'))
model.add(layers.MaxPooling2D((2,2)))
model.add(layers.Flatten())
model.add(layers.Dropout(.5))
model.add(layers.Dense(512, activation='relu',
                        kernel_initializer = 'glorot_normal',
                        bias_initializer = 'zeros'))
model.add(layers.Dense(512, activation='relu'))
model.add(layers.Dense(5, activation='softmax'))
model.summary()
```

Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|---|----------------------|---------|
| ===== | | |
| conv2d_4 (Conv2D) | (None, 126, 126, 32) | 896 |
| batch_normalization_1 (Batch Normalization) | (None, 126, 126, 32) | 128 |
| max_pooling2d_4 (MaxPooling2D) | (None, 63, 63, 32) | 0 |
| conv2d_5 (Conv2D) | (None, 61, 61, 64) | 18496 |
| max_pooling2d_5 (MaxPooling2D) | (None, 30, 30, 64) | 0 |
| conv2d_6 (Conv2D) | (None, 28, 28, 128) | 73856 |
| max_pooling2d_6 (MaxPooling2D) | (None, 14, 14, 128) | 0 |
| conv2d_7 (Conv2D) | (None, 12, 12, 128) | 147584 |
| max_pooling2d_7 (MaxPooling2D) | (None, 6, 6, 128) | 0 |
| flatten_1 (Flatten) | (None, 4608) | 0 |
| dropout_1 (Dropout) | (None, 4608) | 0 |

| | | |
|-----------------------------|-------------|---------|
| dense_3 (Dense) | (None, 512) | 2359808 |
| dense_4 (Dense) | (None, 512) | 262656 |
| dense_5 (Dense) | (None, 5) | 2565 |
| ===== | | |
| Total params: 2,865,989 | | |
| Trainable params: 2,865,925 | | |
| Non-trainable params: 64 | | |

```
In [15]: adam = optimizers.Adam(lr = 0.001)

model.compile(loss = "categorical_crossentropy",
              optimizer=adam,
              metrics=["acc"])
```

In [16]: batch_size=32

Training model with adagrad, 30 epochs and shuffling data

```
history_adam = model.fit_generator(train_generator,
                                   steps_per_epoch=int(444/batch_size),
                                   epochs=30,
                                   validation_data=validation_generator,
                                   validation_steps=int(110/batch_size),
                                   shuffle=True)
```

WARNING:tensorflow:sample_weight modes were coerced from

```
...
to
['...']
```

WARNING:tensorflow:sample_weight modes were coerced from

```
...
to
['...']
```

Train for 13 steps, validate for 3 steps

Epoch 1/30

13/13 [=====] - 21s 2s/step - loss: 1.9139 - acc: 0.1699 - val_loss: 1.6111 - val_acc: 0.0833

Epoch 2/30

13/13 [=====] - 17s 1s/step - loss: 1.6174 - acc: 0.2063 - val_loss: 1.6122 - val_acc: 0.2083

Epoch 3/30

13/13 [=====] - 17s 1s/step - loss: 1.6096 - acc: 0.2354 - val_loss: 1.6123 - val_acc: 0.0833

Epoch 4/30

13/13 [=====] - 18s 1s/step - loss: 1.6059 - acc: 0.2403 - val_loss: 1.6095 - val_acc: 0.1042

Epoch 5/30

13/13 [=====] - 17s 1s/step - loss: 1.5978 - acc: 0.2451 - val_loss: 1.6012 - val_acc: 0.2917

Epoch 6/30

13/13 [=====] - 17s 1s/step - loss: 1.5891 - acc: 0.2597 - val_loss: 1.5727 - val_acc: 0.2396

Epoch 7/30

13/13 [=====] - 17s 1s/step - loss: 1.5578 - acc: 0.3107 - val_loss: 1.5493 - val_acc: 0.4271

Epoch 8/30

13/13 [=====] - 17s 1s/step - loss: 1.5622 - acc: 0.2933 - val_loss: 1.5636 - val_acc: 0.3646

Epoch 9/30

13/13 [=====] - 17s 1s/step - loss: 1.5144 - acc: 0.3131 - val_loss: 1.5499 - val_acc: 0.2604

Epoch 10/30

13/13 [=====] - 10s 752ms/step - loss: 1.5888 - acc: 0.2500 - val_loss: 1.5813 - val_acc: 0.2500

Epoch 11/30

13/13 [=====] - 8s 578ms/step - loss: 1.5253 - acc: 0.3131 - val_loss: 1.5848 - val_acc: 0.1875

Epoch 12/30

13/13 [=====] - 10s 743ms/step - loss: 1.4866 - acc: 0.3495 - val_loss: 1.4979 - val_acc: 0.3229

Epoch 13/30

13/13 [=====] - 10s 733ms/step - loss: 1.4982 - acc: 0.3665 - val_loss: 1.4878 - val_acc: 0.3438

Epoch 14/30

13/13 [=====] - 8s 652ms/step - loss: 1.4756 - acc: 0.3301 -
val_loss: 1.4636 - val_acc: 0.3542
Epoch 15/30
13/13 [=====] - 8s 613ms/step - loss: 1.3970 - acc: 0.3519 -
val_loss: 1.5178 - val_acc: 0.2500
Epoch 16/30
13/13 [=====] - 9s 682ms/step - loss: 1.4459 - acc: 0.3592 -
val_loss: 1.4380 - val_acc: 0.3646
Epoch 17/30
13/13 [=====] - 10s 797ms/step - loss: 1.4367 - acc: 0.3374 -
val_loss: 1.4100 - val_acc: 0.3646
Epoch 18/30
13/13 [=====] - 9s 679ms/step - loss: 1.3735 - acc: 0.3762 -
val_loss: 1.5018 - val_acc: 0.2708
Epoch 19/30
13/13 [=====] - 8s 604ms/step - loss: 1.4102 - acc: 0.3519 -
val_loss: 1.5192 - val_acc: 0.2500
Epoch 20/30
13/13 [=====] - 8s 621ms/step - loss: 1.3858 - acc: 0.3811 -
val_loss: 1.3685 - val_acc: 0.4167
Epoch 21/30
13/13 [=====] - 9s 669ms/step - loss: 1.3434 - acc: 0.3665 -
val_loss: 1.5061 - val_acc: 0.2917
Epoch 22/30
13/13 [=====] - 8s 611ms/step - loss: 1.3317 - acc: 0.3956 -
val_loss: 1.2839 - val_acc: 0.4062
Epoch 23/30
13/13 [=====] - 9s 665ms/step - loss: 1.2943 - acc: 0.4345 -
val_loss: 1.3253 - val_acc: 0.3438
Epoch 24/30
13/13 [=====] - 9s 682ms/step - loss: 1.3110 - acc: 0.4126 -
val_loss: 1.4418 - val_acc: 0.3854
Epoch 25/30
13/13 [=====] - 8s 634ms/step - loss: 1.2613 - acc: 0.4029 -
val_loss: 1.2648 - val_acc: 0.3854
Epoch 26/30
13/13 [=====] - 9s 730ms/step - loss: 1.2704 - acc: 0.4442 -
val_loss: 1.2649 - val_acc: 0.4271
Epoch 27/30
13/13 [=====] - 9s 710ms/step - loss: 1.2252 - acc: 0.4393 -
val_loss: 1.1026 - val_acc: 0.4479
Epoch 28/30
13/13 [=====] - 8s 636ms/step - loss: 1.2611 - acc: 0.4709 -
val_loss: 1.3259 - val_acc: 0.3646
Epoch 29/30
13/13 [=====] - 8s 616ms/step - loss: 1.1703 - acc: 0.4490 -
val_loss: 1.1719 - val_acc: 0.4583
Epoch 30/30
13/13 [=====] - 9s 654ms/step - loss: 1.1776 - acc: 0.4976 -
val_loss: 1.1865 - val_acc: 0.4583

```

In [17]: # -----
# Visualizing Train/Validation Loss & Accuracy /
# -----

acc_adam = history_adam.history['acc']
val_acc_adam = history_adam.history['val_acc']
loss_adam = history_adam.history['loss']
val_loss_adam = history_adam.history['val_loss']

epochs_adam = range(1, len(acc_adam) + 1)


# Plot of accuracy
plt.plot(epochs_adam, acc_adam, color='blue', label='Training Acc')
plt.plot(epochs_adam, val_acc_adam, color='red', label='Validation Acc')
plt.title('Training and Validation Accuracy (Adam)')
plt.legend()


plt.figure()

# Plot of loss
plt.plot(epochs_adam, loss_adam, color='blue', label='Training Loss')
plt.plot(epochs_adam, val_loss_adam, color='red', label='Validation Loss')
plt.title('Training and Validation Loss (Adam)')
plt.legend()
plt.figure()

plt.show()

```





<Figure size 432x288 with 0 Axes>

```
In [18]: # -----
# SAVE THE MODEL |
# -----
model.save('DIP_Proj_modelAdam.h5')
```

```
In [24]: # -----
# Confusion Matrix |
# -----

num_of_train_samples = 444
#num_of_test_samples = 416 # steps per epoch
batch_size=32
steps_per_epoch=num_of_train_samples // batch_size
#validation_generator.reset()

Y_pred = model.predict_generator(validation_generator, steps_per_epoch)
y_pred = np.argmax(Y_pred, axis=1)
print('Confusion Matrix for Model with Adam Optimizer')
print(confusion_matrix(validation_generator.classes, y_pred))

-----
AttributeError                                Traceback (most recent call last)
<ipython-input-24-2188fe66f14c> in <module>
      9 #validation_generator.reset()
     10
--> 11 Y_pred = model.predict_generator(validation_generator, steps_per_epoch)
     12 y_pred = np.argmax(Y_pred, axis=1)
     13 print('Confusion Matrix for Model with Adam Optimizer')

AttributeError: 'Sequential' object has no attribute 'predict_generator'
```

```
In [25]: # -----
# Classification Report /
# -----

print('Classification Report for Model with Adam Optimizer')
target_names = ['1_Neutral', '2_Smiling', '3_Sleepy', '4_Surprise', '5_Sunglasses']
print(classification_report(validation_generator.classes, y_pred, target_names=target_names))
```

```
Classification Report for Model with Adam Optimizer
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 1_Neutral | 0.36 | 0.64 | 0.46 | 22 |
| 2_Smiling | 0.00 | 0.00 | 0.00 | 22 |
| 3_Sleepy | 0.00 | 0.00 | 0.00 | 22 |
| 4_Surprise | 0.43 | 1.00 | 0.60 | 22 |
| 5_Sunglasses | 1.00 | 0.91 | 0.95 | 22 |
| accuracy | | | 0.51 | 110 |
| macro avg | 0.36 | 0.51 | 0.40 | 110 |
| weighted avg | 0.36 | 0.51 | 0.40 | 110 |

C:\Users\User\anaconda3\envs\tensorflow\lib\site-packages\sklearn\metrics_classification.py:1221: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```



```

In [22]: # -----
# ROC/AUC Score |
# -----

from sklearn.preprocessing import LabelBinarizer

# set plot figure size
fig, c_ax = plt.subplots(1,1, figsize = (12, 8))

def multiclass_roc_auc_score(y_test, y_pred, average="macro"):
    lb = LabelBinarizer()
    lb.fit(y_test)
    y_test = lb.transform(y_test)
    y_pred = lb.transform(y_pred)

    for (idx, c_label) in enumerate(target_names): # target_names: no of the labels
        fpr, tpr, thresholds = roc_curve(y_test[:,idx].astype(int), y_pred[:,idx])
        c_ax.plot(fpr, tpr, label = '%s (AUC:%0.2f)' % (c_label, auc(fpr, tpr)))
    c_ax.plot(fpr, fpr, 'b-', label = 'Random Guessing')
    return roc_auc_score(y_test, y_pred, average=average)

validation_generator.reset() # resetting generator
y_pred = model.predict_generator(validation_generator, verbose = True)
y_pred = np.argmax(y_pred, axis=1)
multiclass_roc_auc_score(validation_generator.classes, y_pred)

```


WARNING:tensorflow:From <ipython-input-22-470a0699b29d>:23: Model.predict_generator (from tensorflow.python.keras.engine.training) is deprecated and will be removed in a future version.

Instructions for updating:

Please use Model.predict, which supports generators.

4/4 [=====] - 0s 107ms/step

Out[22]: 0.6931818181818181



```

In [28]: # -----
# Displaying 12 images with the prediction |
# -----

# 1 image from each class

labels = ['1_Neutral', '2_Smiling', '3_Sleepy', '4_Surprise', '5_Sunglasses']


i=0
for names in labels:
    pathToFolder = test_dir + '/' + names + '/'
    fnames = [os.path.join(pathToFolder, fname) for
               fname in os.listdir(pathToFolder)]

    # generate random number btwn (0,30)
    randNum = random.randint(0, 20)
    img_path = fnames[randNum]
    tmp_img = image.load_img(img_path, target_size = (128, 128))
    tmp_img = image.img_to_array(tmp_img)
    tmp_img = np.expand_dims(tmp_img, axis = 0)


    tmp_img /=255.
    plt.imshow(tmp_img[0])
    plt.show()
    # predict
    result = model.predict(tmp_img)
    train_generator.class_indices
    print('Actual value: ', i, '\tPredicted value: ', np.argmax(result))
    i+=1

print('Total number of images for "testing":')
test_generator = test_datagen.flow_from_directory(test_dir,
                                                    target_size = (128, 128),
                                                    batch_size = 32,
                                                    class_mode = "categorical",
                                                    shuffle=False)


```




Actual value: 0 Predicted value: 0




Actual value: 1 Predicted value: 3



Actual value: 2 Predicted value: 3



Actual value: 3 Predicted value: 3



Actual value: 4 Predicted value: 4
Total number of images for "testing":
Found 110 images belonging to 5 classes.

```
In [186]: # MODEL 2
# -----
# ReLU activation function
# Adadagrad optimizer
# 4 Conv
# 1 Batch Norm
# 4 MaxPool
# 1 Dropout
# 3 Dense
# 1 Flatten
# Batchsize 32

model2 = tensorflow.keras.Sequential()
model2.add(layers.Conv2D(32, (3,3), activation='relu',
                        input_shape = (128, 128, 3),
                        kernel_initializer = 'glorot_normal',
                        bias_initializer = 'zeros'))
model2.add(layers.BatchNormalization())
model2.add(layers.MaxPooling2D((2,2)))
model2.add(layers.Conv2D(64, (3,3), activation='relu'))
model2.add(layers.MaxPooling2D((2,2)))
model2.add(layers.Conv2D(128, (3,3), activation='relu'))
model2.add(layers.MaxPooling2D((2,2)))
model2.add(layers.Conv2D(128, (3,3), activation='relu'))
model2.add(layers.MaxPooling2D((2,2)))
model2.add(layers.Flatten())
model2.add(layers.Dropout(.5))
model2.add(layers.Dense(32, activation='relu',
                        kernel_initializer = 'glorot_normal',
                        bias_initializer = 'zeros'))
model2.add(layers.Dense(64, activation='relu'))
model2.add(layers.Dense(5, activation='softmax'))
model2.summary()
```

Model: "sequential_39"

| Layer (type) | Output Shape | Param # |
|--|----------------------|---------|
| ===== | | |
| conv2d_155 (Conv2D) | (None, 126, 126, 32) | 896 |
| batch_normalization_39 (Batch Normalization) | (None, 126, 126, 32) | 128 |
| max_pooling2d_155 (MaxPooling2D) | (None, 63, 63, 32) | 0 |
| conv2d_156 (Conv2D) | (None, 61, 61, 64) | 18496 |
| max_pooling2d_156 (MaxPooling2D) | (None, 30, 30, 64) | 0 |
| conv2d_157 (Conv2D) | (None, 28, 28, 128) | 73856 |
| max_pooling2d_157 (MaxPooling2D) | (None, 14, 14, 128) | 0 |
| conv2d_158 (Conv2D) | (None, 12, 12, 128) | 147584 |
| max_pooling2d_158 (MaxPooling2D) | (None, 6, 6, 128) | 0 |
| flatten_39 (Flatten) | (None, 4608) | 0 |
| dropout_39 (Dropout) | (None, 4608) | 0 |

| | | |
|---------------------------|------------|--------|
| dense_115 (Dense) | (None, 32) | 147488 |
| dense_116 (Dense) | (None, 64) | 2112 |
| dense_117 (Dense) | (None, 5) | 325 |
| ===== | | |
| Total params: 390,885 | | |
| Trainable params: 390,821 | | |
| Non-trainable params: 64 | | |

```
In [187]: adagrad = optimizers.Adagrad(lr=0.01)

model2.compile(loss = "categorical_crossentropy",
               optimizer=adagrad,
               metrics=["acc"])
```

In [188]: batch_size=32

Training model with adagrad, 30 epochs and shuffling data

```
history_adagrad = model2.fit_generator(train_generator,
                                       steps_per_epoch=int(444/batch_size),
                                       epochs=50,
                                       validation_data=validation_generator,
                                       validation_steps=int(110/batch_size),
                                       shuffle=True)
```

WARNING:tensorflow:sample_weight modes were coerced from

```
...
to
['...']
```

WARNING:tensorflow:sample_weight modes were coerced from

```
...
to
['...']
```

Train for 13 steps, validate for 3 steps

Epoch 1/50

13/13 [=====] - 8s 628ms/step - loss: 1.6666 - acc: 0.2015 -
val_loss: 1.6109 - val_acc: 0.2292

Epoch 2/50

13/13 [=====] - 8s 592ms/step - loss: 1.6132 - acc: 0.1966 -
val_loss: 1.6129 - val_acc: 0.0833

Epoch 3/50

13/13 [=====] - 8s 597ms/step - loss: 1.6147 - acc: 0.2282 -
val_loss: 1.6059 - val_acc: 0.2396

Epoch 4/50

13/13 [=====] - 8s 613ms/step - loss: 1.6131 - acc: 0.2160 -
val_loss: 1.6097 - val_acc: 0.0833

Epoch 5/50

13/13 [=====] - 8s 589ms/step - loss: 1.6090 - acc: 0.2209 -
val_loss: 1.6070 - val_acc: 0.2708

Epoch 6/50

13/13 [=====] - 8s 598ms/step - loss: 1.6145 - acc: 0.1917 -
val_loss: 1.6074 - val_acc: 0.2188

Epoch 7/50

13/13 [=====] - 8s 596ms/step - loss: 1.6180 - acc: 0.1748 -
val_loss: 1.6061 - val_acc: 0.2292

Epoch 8/50

13/13 [=====] - 8s 604ms/step - loss: 1.6068 - acc: 0.2476 -
val_loss: 1.6021 - val_acc: 0.2083

Epoch 9/50

13/13 [=====] - 8s 593ms/step - loss: 1.6059 - acc: 0.2039 -
val_loss: 1.6044 - val_acc: 0.3333

Epoch 10/50

13/13 [=====] - 8s 598ms/step - loss: 1.6087 - acc: 0.2354 -
val_loss: 1.6030 - val_acc: 0.3229

Epoch 11/50

13/13 [=====] - 8s 598ms/step - loss: 1.5897 - acc: 0.2427 -
val_loss: 1.5974 - val_acc: 0.2396

Epoch 12/50

13/13 [=====] - 8s 590ms/step - loss: 1.5979 - acc: 0.2160 -
val_loss: 1.5950 - val_acc: 0.3229

Epoch 13/50

13/13 [=====] - 8s 591ms/step - loss: 1.6020 - acc: 0.2257 -
val_loss: 1.5932 - val_acc: 0.2812


Epoch 14/50


13/13 [=====] - 8s 593ms/step - loss: 1.6033 - acc: 0.2112 -
val_loss: 1.5948 - val_acc: 0.2604
Epoch 15/50
13/13 [=====] - 8s 594ms/step - loss: 1.5769 - acc: 0.2670 -
val_loss: 1.5721 - val_acc: 0.2396
Epoch 16/50
13/13 [=====] - 8s 596ms/step - loss: 1.5703 - acc: 0.2549 -
val_loss: 1.5817 - val_acc: 0.2500
Epoch 17/50
13/13 [=====] - 8s 602ms/step - loss: 1.5699 - acc: 0.2791 -
val_loss: 1.5562 - val_acc: 0.3438
Epoch 18/50
13/13 [=====] - 8s 594ms/step - loss: 1.5495 - acc: 0.2937 -
val_loss: 1.5440 - val_acc: 0.2812
Epoch 19/50
13/13 [=====] - 8s 591ms/step - loss: 1.5701 - acc: 0.2743 -
val_loss: 1.5350 - val_acc: 0.3333
Epoch 20/50
13/13 [=====] - 8s 597ms/step - loss: 1.5237 - acc: 0.3228 -
val_loss: 1.5037 - val_acc: 0.3229
Epoch 21/50
13/13 [=====] - 8s 594ms/step - loss: 1.5486 - acc: 0.2985 -
val_loss: 1.4859 - val_acc: 0.3438
Epoch 22/50
13/13 [=====] - 8s 600ms/step - loss: 1.5433 - acc: 0.2985 -
val_loss: 1.5188 - val_acc: 0.2812
Epoch 23/50
13/13 [=====] - 8s 600ms/step - loss: 1.5277 - acc: 0.2913 -
val_loss: 1.4930 - val_acc: 0.3438
Epoch 24/50
13/13 [=====] - 8s 594ms/step - loss: 1.5138 - acc: 0.3058 -
val_loss: 1.4770 - val_acc: 0.2708
Epoch 25/50
13/13 [=====] - 8s 611ms/step - loss: 1.5027 - acc: 0.3173 -
val_loss: 1.4419 - val_acc: 0.3125
Epoch 26/50
13/13 [=====] - 8s 631ms/step - loss: 1.5028 - acc: 0.3350 -
val_loss: 1.4398 - val_acc: 0.3542
Epoch 27/50
13/13 [=====] - 8s 608ms/step - loss: 1.5241 - acc: 0.3277 -
val_loss: 1.4360 - val_acc: 0.3646
Epoch 28/50
13/13 [=====] - 8s 590ms/step - loss: 1.4486 - acc: 0.3665 -
val_loss: 1.3718 - val_acc: 0.4375
Epoch 29/50
13/13 [=====] - 8s 593ms/step - loss: 1.4553 - acc: 0.3738 -
val_loss: 1.3754 - val_acc: 0.3646
Epoch 30/50
13/13 [=====] - 8s 586ms/step - loss: 1.4732 - acc: 0.3617 -
val_loss: 1.3863 - val_acc: 0.4062
Epoch 31/50
13/13 [=====] - 8s 594ms/step - loss: 1.4567 - acc: 0.3252 -
val_loss: 1.4100 - val_acc: 0.3542
Epoch 32/50
13/13 [=====] - 8s 589ms/step - loss: 1.4424 - acc: 0.3811 -
val_loss: 1.3363 - val_acc: 0.3958
Epoch 33/50
13/13 [=====] - 8s 594ms/step - loss: 1.4560 - acc: 0.3495 -
val_loss: 1.3591 - val_acc: 0.3542

Epoch 34/50
13/13 [=====] - 8s 603ms/step - loss: 1.4350 - acc: 0.3859 -
val_loss: 1.3509 - val_acc: 0.4479
Epoch 35/50
13/13 [=====] - 8s 603ms/step - loss: 1.4361 - acc: 0.3665 -
val_loss: 1.3002 - val_acc: 0.4375
Epoch 36/50
13/13 [=====] - 8s 598ms/step - loss: 1.3867 - acc: 0.4053 -
val_loss: 1.2712 - val_acc: 0.4688
Epoch 37/50
13/13 [=====] - 8s 597ms/step - loss: 1.4081 - acc: 0.3495 -
val_loss: 1.4801 - val_acc: 0.3333
Epoch 38/50
13/13 [=====] - 8s 601ms/step - loss: 1.3693 - acc: 0.3981 -
val_loss: 1.2789 - val_acc: 0.5521
Epoch 39/50
13/13 [=====] - 8s 606ms/step - loss: 1.4014 - acc: 0.3689 -
val_loss: 1.2286 - val_acc: 0.3854
Epoch 40/50
13/13 [=====] - 9s 685ms/step - loss: 1.3862 - acc: 0.3859 -
val_loss: 1.1689 - val_acc: 0.4583
Epoch 41/50
13/13 [=====] - 8s 602ms/step - loss: 1.3711 - acc: 0.4087 -
val_loss: 1.1800 - val_acc: 0.4375
Epoch 42/50
13/13 [=====] - 8s 614ms/step - loss: 1.3349 - acc: 0.4062 -
val_loss: 1.1561 - val_acc: 0.5000
Epoch 43/50
13/13 [=====] - 10s 760ms/step - loss: 1.3782 - acc: 0.3908 -
val_loss: 1.1448 - val_acc: 0.4896
Epoch 44/50
13/13 [=====] - 9s 690ms/step - loss: 1.3409 - acc: 0.4078 -
val_loss: 1.1163 - val_acc: 0.5000
Epoch 45/50
13/13 [=====] - 8s 598ms/step - loss: 1.3373 - acc: 0.4150 -
val_loss: 1.1115 - val_acc: 0.5729
Epoch 46/50
13/13 [=====] - 8s 601ms/step - loss: 1.3310 - acc: 0.4199 -
val_loss: 1.1506 - val_acc: 0.4896
Epoch 47/50
13/13 [=====] - 8s 640ms/step - loss: 1.3069 - acc: 0.4248 -
val_loss: 1.0899 - val_acc: 0.6042
Epoch 48/50
13/13 [=====] - 8s 601ms/step - loss: 1.3100 - acc: 0.4296 -
val_loss: 1.1005 - val_acc: 0.5625
Epoch 49/50
13/13 [=====] - 8s 609ms/step - loss: 1.2631 - acc: 0.4490 -
val_loss: 1.1794 - val_acc: 0.4375
Epoch 50/50
13/13 [=====] - 8s 647ms/step - loss: 1.3043 - acc: 0.4442 -
val_loss: 1.1197 - val_acc: 0.5312

In [189]:

```
# -----  
# Visualizing Train/Validation Loss & Accuracy /  
# -----  
  
acc_adagrad = history_adagrad.history['acc']  
val_acc_adagrad = history_adagrad.history['val_acc']  
loss_adagrad = history_adagrad.history['loss']  
val_loss_adagrad = history_adagrad.history['val_loss']  
  
epochs_adagrad = range(1, len(acc_adagrad) + 1)  
  
# Plot of accuracy  
plt.plot(epochs_adagrad, acc_adagrad, color='blue', label='Training Acc')  
plt.plot(epochs_adagrad, val_acc_adagrad, color='red', label='Validation Acc')  
plt.title('Training and Validation Accuracy (Adagrad)')  
plt.legend()  
  
plt.figure()  
  
# Plot of loss  
plt.plot(epochs_adagrad, loss_adagrad, color='blue', label='Training Loss')  
plt.plot(epochs_adagrad, val_loss_adagrad, color='red', label='Validation Loss')  
plt.title('Training and Validation Loss (Adagrad)')  
plt.legend()  
plt.figure()  
  
plt.show()
```





<Figure size 432x288 with 0 Axes>

```

In [190]: # -----
# ROC/AUC Score |
# -----
from sklearn.preprocessing import LabelBinarizer

# set plot figure size
fig, c_ax = plt.subplots(1,1, figsize = (12, 8))

def multiclass_roc_auc_score(y_test, y_pred, average="macro"):
    lb = LabelBinarizer()
    lb.fit(y_test)
    y_test = lb.transform(y_test)
    y_pred = lb.transform(y_pred)


    for (idx, c_label) in enumerate(target_names): # target_names: no of the labels
        fpr, tpr, thresholds = roc_curve(y_test[:,idx].astype(int), y_pred[:,idx])
        c_ax.plot(fpr, tpr, label = '%s (AUC:%0.2f)' % (c_label, auc(fpr, tpr)))
    c_ax.plot(fpr, fpr, 'b-', label = 'Random Guessing')
    return roc_auc_score(y_test, y_pred, average=average)

validation_generator.reset() # resetting generator
y_pred = model2.predict_generator(validation_generator, verbose = True)
y_pred = np.argmax(y_pred, axis=1)
multiclass_roc_auc_score(validation_generator.classes, y_pred)

```

4/4 [=====] - 0s 98ms/step

Out[190]: 0.7443181818181819



```
In [191]: # -----
# Classification Report /
# -----

print('Classification Report for Model with Adagrad')
target_names = ['1_Neutral', '2_Smiling', '3_Sleepy', '4_Surprise', '5_Sunglasses']
print(classification_report(validation_generator.classes, y_pred, target_names=target_names))
```

```
Classification Report for Model with Adagrad
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 1_Neutral | 0.45 | 0.23 | 0.30 | 22 |
| 2_Smiling | 0.33 | 0.09 | 0.14 | 22 |
| 3_Sleepy | 0.37 | 0.77 | 0.50 | 22 |
| 4_Surprise | 0.90 | 0.86 | 0.88 | 22 |
| 5_Sunglasses | 0.85 | 1.00 | 0.92 | 22 |
| accuracy | | | 0.59 | 110 |
| macro avg | 0.58 | 0.59 | 0.55 | 110 |
| weighted avg | 0.58 | 0.59 | 0.55 | 110 |

```
In [192]: # -----
# Confusion Matrix /
# -----

num_of_train_samples = 444
#num_of_test_samples = 416 # steps per epoch
batch_size=32
steps_per_epoch=num_of_train_samples // batch_size
#validation_generator.reset()

Y_pred = model2.predict_generator(validation_generator, steps_per_epoch)
y_pred = np.argmax(Y_pred, axis=1)
print('Confusion Matrix fo Model with Adagrad')
print(confusion_matrix(validation_generator.classes, y_pred))
```

WARNING:tensorflow:Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches (in this case, 13 batches). You may need to use the repeat() function when building your dataset.

Confusion Matrix fo Model with Adagrad

```
[[ 5  2 14  1  0]
 [ 3  2 14  1  2]
 [ 2  1 17  0  2]
 [ 1  1  1 19  0]
 [ 0  0  0  0 22]]
```

```

In [193]: # -----
# Displaying 12 Images with the Prediction /
# -----

# 1 image from each class

labels = ['1_Neutral', '2_Smiling', '3_Sleepy', '4_Surprise', '5_Sunglasses']


i=0
for names in labels:
    pathToFolder = test_dir + '/' + names + '/'
    fnames = [os.path.join(pathToFolder, fname) for
               fname in os.listdir(pathToFolder)]

    # generate random number btwn (0,30)
    randNum = random.randint(0, 20)
    img_path = fnames[randNum]
    tmp_img = image.load_img(img_path, target_size = (128, 128))
    tmp_img = image.img_to_array(tmp_img)
    tmp_img = np.expand_dims(tmp_img, axis = 0)


    tmp_img /=255.
    plt.imshow(tmp_img[0])
    plt.show()
    # predict
    result = model2.predict(tmp_img)
    train_generator.class_indices
    print('Actual value: ',i,'\tPredicted value: ', np.argmax(result))
    i+=1

print('Total number of images for "testing":')
test_generator = test_datagen.flow_from_directory(test_dir,
                                                    target_size = (128, 128),
                                                    batch_size = 32,
                                                    class_mode = "categorical",
                                                    shuffle=False)


```




Actual value: 0 Predicted value: 2



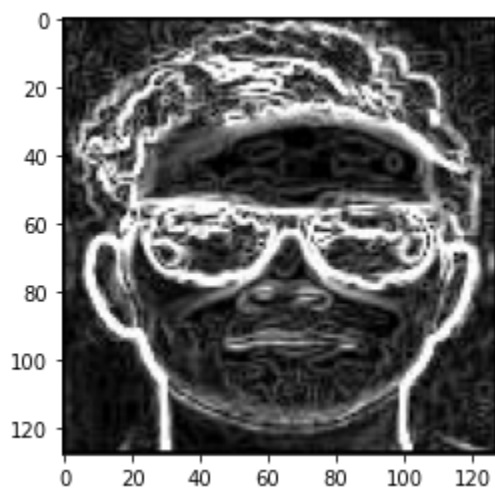
Actual value: 1 Predicted value: 0



Actual value: 2 Predicted value: 0



Actual value: 3 Predicted value: 2



Actual value: 4 Predicted value: 4
Total number of images for "testing":
Found 110 images belonging to 5 classes.

```
In [133]: # MODEL 3
# -----
# ReLU activation function
# Rmsprop optimizer
# 4 Conv
# 1 Batch Norm
# 4 MaxPool
# 1 Dropout
# 3 Dense
# 1 Flatten
# Batchsize 32

model3 = tensorflow.keras.Sequential()
model3.add(layers.Conv2D(32, (3,3), activation='relu',
                        input_shape = (128, 128, 3),
                        kernel_initializer = 'glorot_normal',
                        bias_initializer = 'zeros'))
model3.add(layers.BatchNormalization())
model3.add(layers.MaxPooling2D((2,2)))
model3.add(layers.Conv2D(64, (3,3), activation='relu'))
model3.add(layers.MaxPooling2D((2,2)))
model3.add(layers.Conv2D(128, (3,3), activation='relu'))
model3.add(layers.MaxPooling2D((2,2)))
model3.add(layers.Conv2D(128, (3,3), activation='relu'))
model3.add(layers.MaxPooling2D((2,2)))
model3.add(layers.Flatten())
model3.add(layers.Dropout(.5))
model3.add(layers.Dense(64, activation='relu',
                        kernel_initializer = 'glorot_normal',
                        bias_initializer = 'zeros'))
model3.add(layers.Dense(128, activation='relu'))
model3.add(layers.Dense(5, activation='softmax'))
model3.summary()
```

Model: "sequential_23"

| Layer (type) | Output Shape | Param # |
|--|----------------------|---------|
| ===== | | |
| conv2d_91 (Conv2D) | (None, 126, 126, 32) | 896 |
| batch_normalization_23 (Batch Normalization) | (None, 126, 126, 32) | 128 |
| max_pooling2d_91 (MaxPooling2D) | (None, 63, 63, 32) | 0 |
| conv2d_92 (Conv2D) | (None, 61, 61, 64) | 18496 |
| max_pooling2d_92 (MaxPooling2D) | (None, 30, 30, 64) | 0 |
| conv2d_93 (Conv2D) | (None, 28, 28, 128) | 73856 |
| max_pooling2d_93 (MaxPooling2D) | (None, 14, 14, 128) | 0 |
| conv2d_94 (Conv2D) | (None, 12, 12, 128) | 147584 |
| max_pooling2d_94 (MaxPooling2D) | (None, 6, 6, 128) | 0 |
| flatten_23 (Flatten) | (None, 4608) | 0 |
| dropout_23 (Dropout) | (None, 4608) | 0 |

| | | |
|---------------------------|-------------|--------|
| dense_67 (Dense) | (None, 64) | 294976 |
| dense_68 (Dense) | (None, 128) | 8320 |
| dense_69 (Dense) | (None, 5) | 645 |
| ===== | | |
| Total params: 544,901 | | |
| Trainable params: 544,837 | | |
| Non-trainable params: 64 | | |
| ===== | | |

```
In [134]: rmsprop = optimizers.RMSprop(lr=0.001, rho=0.9)

model3.compile(loss = "categorical_crossentropy",
               optimizer=rmsprop,
               metrics=["acc"])
```

In [135]: *# Training model with adagrad, 30 epochs and shuffling data*

```
batch_size=32
```

```
history_rsm = model3.fit_generator(train_generator,  
                                   steps_per_epoch=int(444/batch_size),  
                                   epochs=50,  
                                   validation_data=validation_generator,  
                                   validation_steps=int(110/batch_size),  
                                   shuffle=True)
```

WARNING:tensorflow:sample_weight modes were coerced from

```
...  
to  
['...']
```

WARNING:tensorflow:sample_weight modes were coerced from

```
...  
to  
['...']
```

Train for 13 steps, validate for 3 steps

Epoch 1/50

13/13 [=====] - 9s 694ms/step - loss: 1.6852 - acc: 0.2039 -
val_loss: 1.6122 - val_acc: 0.0833

Epoch 2/50

13/13 [=====] - 8s 634ms/step - loss: 1.6243 - acc: 0.1845 -
val_loss: 1.6109 - val_acc: 0.0833

Epoch 3/50

13/13 [=====] - 9s 690ms/step - loss: 1.6147 - acc: 0.2257 -
val_loss: 1.6082 - val_acc: 0.2292

Epoch 4/50

13/13 [=====] - 8s 594ms/step - loss: 1.6061 - acc: 0.2500 -
val_loss: 1.6060 - val_acc: 0.2708

Epoch 5/50

13/13 [=====] - 8s 596ms/step - loss: 1.5989 - acc: 0.2743 -
val_loss: 1.6158 - val_acc: 0.1042

Epoch 6/50

13/13 [=====] - 8s 606ms/step - loss: 1.5667 - acc: 0.2500 -
val_loss: 1.5626 - val_acc: 0.3021

Epoch 7/50

13/13 [=====] - 8s 599ms/step - loss: 1.5565 - acc: 0.2937 -
val_loss: 1.5526 - val_acc: 0.3125

Epoch 8/50

13/13 [=====] - 8s 599ms/step - loss: 1.5201 - acc: 0.3053 -
val_loss: 1.6186 - val_acc: 0.2083

Epoch 9/50

13/13 [=====] - 8s 596ms/step - loss: 1.5536 - acc: 0.3197 -
val_loss: 1.5098 - val_acc: 0.3646

Epoch 10/50

13/13 [=====] - 8s 607ms/step - loss: 1.4868 - acc: 0.3398 -
val_loss: 1.4654 - val_acc: 0.4062

Epoch 11/50

13/13 [=====] - 8s 593ms/step - loss: 1.4285 - acc: 0.3641 -
val_loss: 1.7033 - val_acc: 0.1458

Epoch 12/50

13/13 [=====] - 8s 602ms/step - loss: 1.5643 - acc: 0.3155 -
val_loss: 1.4140 - val_acc: 0.4167

Epoch 13/50

13/13 [=====] - 8s 591ms/step - loss: 1.4411 - acc: 0.3654 -
val_loss: 1.3967 - val_acc: 0.5104

Epoch 14/50

13/13 [=====] - 8s 590ms/step - loss: 1.4523 - acc: 0.3422 -
val_loss: 1.3572 - val_acc: 0.4271
Epoch 15/50
13/13 [=====] - 8s 609ms/step - loss: 1.4170 - acc: 0.4175 -
val_loss: 1.2582 - val_acc: 0.5208
Epoch 16/50
13/13 [=====] - 8s 598ms/step - loss: 1.3619 - acc: 0.3762 -
val_loss: 1.2577 - val_acc: 0.4271
Epoch 17/50
13/13 [=====] - 8s 591ms/step - loss: 1.3416 - acc: 0.4369 -
val_loss: 1.1675 - val_acc: 0.5104
Epoch 18/50
13/13 [=====] - 8s 591ms/step - loss: 1.3363 - acc: 0.4199 -
val_loss: 1.2245 - val_acc: 0.5104
Epoch 19/50
13/13 [=====] - 8s 600ms/step - loss: 1.2582 - acc: 0.4806 -
val_loss: 1.1004 - val_acc: 0.6042
Epoch 20/50
13/13 [=====] - 8s 596ms/step - loss: 1.2771 - acc: 0.4490 -
val_loss: 1.1394 - val_acc: 0.4896
Epoch 21/50
13/13 [=====] - 8s 598ms/step - loss: 1.2949 - acc: 0.4369 -
val_loss: 1.0682 - val_acc: 0.5938
Epoch 22/50
13/13 [=====] - 8s 596ms/step - loss: 1.2143 - acc: 0.4757 -
val_loss: 1.0315 - val_acc: 0.5833
Epoch 23/50
13/13 [=====] - 8s 611ms/step - loss: 1.2302 - acc: 0.4515 -
val_loss: 0.9393 - val_acc: 0.6771
Epoch 24/50
13/13 [=====] - 8s 594ms/step - loss: 1.1804 - acc: 0.5097 -
val_loss: 0.9813 - val_acc: 0.6771
Epoch 25/50
13/13 [=====] - 8s 596ms/step - loss: 1.0919 - acc: 0.5194 -
val_loss: 0.9657 - val_acc: 0.5417
Epoch 26/50
13/13 [=====] - 8s 600ms/step - loss: 1.1032 - acc: 0.5243 -
val_loss: 0.8919 - val_acc: 0.7188
Epoch 27/50
13/13 [=====] - 8s 620ms/step - loss: 1.0707 - acc: 0.5388 -
val_loss: 0.8532 - val_acc: 0.7396
Epoch 28/50
13/13 [=====] - 8s 597ms/step - loss: 1.1345 - acc: 0.5607 -
val_loss: 0.8259 - val_acc: 0.7083
Epoch 29/50
13/13 [=====] - 8s 600ms/step - loss: 1.0928 - acc: 0.5364 -
val_loss: 0.8089 - val_acc: 0.6771
Epoch 30/50
13/13 [=====] - 8s 601ms/step - loss: 1.0828 - acc: 0.5769 -
val_loss: 0.7362 - val_acc: 0.7083
Epoch 31/50
13/13 [=====] - 8s 599ms/step - loss: 1.0360 - acc: 0.5461 -
val_loss: 0.7606 - val_acc: 0.6771
Epoch 32/50
13/13 [=====] - 8s 609ms/step - loss: 0.9954 - acc: 0.5655 -
val_loss: 0.6879 - val_acc: 0.6979
Epoch 33/50
13/13 [=====] - 8s 592ms/step - loss: 0.9260 - acc: 0.6019 -
val_loss: 0.6352 - val_acc: 0.7604

```

Epoch 34/50
13/13 [=====] - 8s 596ms/step - loss: 0.9456 - acc: 0.6117 -
val_loss: 0.6991 - val_acc: 0.7292
Epoch 35/50
13/13 [=====] - 8s 598ms/step - loss: 0.9411 - acc: 0.6189 -
val_loss: 0.6508 - val_acc: 0.6771
Epoch 36/50
13/13 [=====] - 8s 612ms/step - loss: 0.9424 - acc: 0.6068 -
val_loss: 0.5583 - val_acc: 0.7604
Epoch 37/50
13/13 [=====] - 8s 639ms/step - loss: 0.9586 - acc: 0.5995 -
val_loss: 0.6296 - val_acc: 0.7604
Epoch 38/50
13/13 [=====] - 8s 634ms/step - loss: 0.8114 - acc: 0.6626 -
val_loss: 0.5484 - val_acc: 0.7812
Epoch 39/50
13/13 [=====] - 8s 599ms/step - loss: 0.8081 - acc: 0.6456 -
val_loss: 0.5287 - val_acc: 0.7083
Epoch 40/50
13/13 [=====] - 8s 610ms/step - loss: 0.8846 - acc: 0.6553 -
val_loss: 0.6120 - val_acc: 0.6979
Epoch 41/50
13/13 [=====] - 8s 592ms/step - loss: 0.7966 - acc: 0.6748 -
val_loss: 0.5689 - val_acc: 0.7604
Epoch 42/50
13/13 [=====] - 8s 592ms/step - loss: 0.7980 - acc: 0.6432 -
val_loss: 0.5261 - val_acc: 0.7188
Epoch 43/50
13/13 [=====] - 8s 595ms/step - loss: 0.8226 - acc: 0.6578 -
val_loss: 0.6015 - val_acc: 0.8021
Epoch 44/50
13/13 [=====] - 8s 596ms/step - loss: 0.7187 - acc: 0.6650 -
val_loss: 0.5520 - val_acc: 0.8333
Epoch 45/50
13/13 [=====] - 8s 598ms/step - loss: 0.8481 - acc: 0.6456 -
val_loss: 0.5600 - val_acc: 0.7500
Epoch 46/50
13/13 [=====] - 8s 596ms/step - loss: 0.7062 - acc: 0.6990 -
val_loss: 0.4593 - val_acc: 0.7396
Epoch 47/50
13/13 [=====] - 8s 592ms/step - loss: 0.7224 - acc: 0.6845 -
val_loss: 0.4998 - val_acc: 0.7708
Epoch 48/50
13/13 [=====] - 8s 601ms/step - loss: 0.7143 - acc: 0.6899 -
val_loss: 0.5317 - val_acc: 0.7708
Epoch 49/50
13/13 [=====] - 8s 600ms/step - loss: 0.6922 - acc: 0.6869 -
val_loss: 0.5935 - val_acc: 0.7188
Epoch 50/50
13/13 [=====] - 8s 590ms/step - loss: 0.7539 - acc: 0.6723 -
val_loss: 0.4173 - val_acc: 0.8438


```


```

In [40]: # -----
# SAVE THE MODEL |
# -----
model3.save('DIP_Proj_modelRSM.h5')

```

```
In [136]: # -----  
# Visualizing Train/Validation Loss & Accuracy /  
# -----  
  
acc_rsm = history_rsm.history['acc']  
val_acc_rsm = history_rsm.history['val_acc']  
loss_rsm = history_rsm.history['loss']  
val_loss_rsm = history_rsm.history['val_loss']  
  
epochs_rsm = range(1, len(acc_rsm) + 1)  
  
# Plot of accuracy  
plt.plot(epochs_rsm, acc_rsm, color='blue', label='Training Acc')  
plt.plot(epochs_rsm, val_acc_rsm, color='red', label='Validation Acc')  
plt.title('Training and Validation Accuracy (RSM)')  
plt.legend()  
  
plt.figure()  
  
# Plot of loss  
plt.plot(epochs_rsm, loss_rsm, color='blue', label='Training Loss')  
plt.plot(epochs_rsm, val_loss_rsm, color='red', label='Validation Loss')  
plt.title('Training and Validation Loss (RSM)')  
plt.legend()  
plt.figure()  
  
plt.show()
```





<Figure size 432x288 with 0 Axes>


```

In [137]: # -----
# ROC/AUC Score |
# -----

from sklearn.preprocessing import LabelBinarizer

# set plot figure size
fig, c_ax = plt.subplots(1,1, figsize = (12, 8))

def multiclass_roc_auc_score(y_test, y_pred, average="macro"):
    lb = LabelBinarizer()
    lb.fit(y_test)
    y_test = lb.transform(y_test)
    y_pred = lb.transform(y_pred)


    for (idx, c_label) in enumerate(target_names): # target_names: no of the labels
        fpr, tpr, thresholds = roc_curve(y_test[:,idx].astype(int), y_pred[:,idx])
        c_ax.plot(fpr, tpr, label = '%s (AUC:%0.2f)' % (c_label, auc(fpr, tpr)))
    c_ax.plot(fpr, fpr, 'b-', label = 'Random Guessing')
    return roc_auc_score(y_test, y_pred, average=average)

validation_generator.reset() # resetting generator
y_pred = model3.predict_generator(validation_generator, verbose = True)
y_pred = np.argmax(y_pred, axis=1)
multiclass_roc_auc_score(validation_generator.classes, y_pred)

```

4/4 [=====] - 0s 100ms/step

Out[137]: 0.9147727272727273



```
In [138]: # -----
# Classification Report /
# -----

print('Classification Report for Model with RSM')
target_names = ['1_Neutral', '2_Smiling', '3_Sleepy', '4_Surprise', '5_Sunglasses']
print(classification_report(validation_generator.classes, y_pred, target_names=target_names))
```

```
Classification Report for Model with RSM
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 1_Neutral | 0.68 | 0.77 | 0.72 | 22 |
| 2_Smiling | 0.79 | 0.68 | 0.73 | 22 |
| 3_Sleepy | 0.90 | 0.86 | 0.88 | 22 |
| 4_Surprise | 0.96 | 1.00 | 0.98 | 22 |
| 5_Sunglasses | 1.00 | 1.00 | 1.00 | 22 |
| accuracy | | | 0.86 | 110 |
| macro avg | 0.87 | 0.86 | 0.86 | 110 |
| weighted avg | 0.87 | 0.86 | 0.86 | 110 |

```
In [139]: # -----
# Confusion Matrix /
# -----

num_of_train_samples = 444
#num_of_test_samples = 416 # steps per epoch
batch_size=32
steps_per_epoch=num_of_train_samples // batch_size
#validation_generator.reset()

Y_pred = model3.predict_generator(validation_generator, steps_per_epoch)
y_pred = np.argmax(Y_pred, axis=1)
print('Confusion Matrix for Model with RSM')
print(confusion_matrix(validation_generator.classes, y_pred))
```

WARNING:tensorflow:Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches (in this case, 13 batches). You may need to use the repeat() function when building your dataset.

Confusion Matrix for Model with RSM

```
[[17  4  1  0  0]
 [ 5 15  1  1  0]
 [ 3  0 19  0  0]
 [ 0  0  0 22  0]
 [ 0  0  0  0 22]]
```

```

In [140]: # -----
# Displaying 12 Images with the Prediction /
# -----

# 1 image from each class

labels = ['1_Neutral', '2_Smiling', '3_Sleepy', '4_Surprise', '5_Sunglasses']


i=0
for names in labels:
    pathToFolder = test_dir + '/' + names + '/'
    fnames = [os.path.join(pathToFolder, fname) for
               fname in os.listdir(pathToFolder)]

    # generate random number btwn (0,30)
    randNum = random.randint(0, 20)
    img_path = fnames[randNum]
    tmp_img = image.load_img(img_path, target_size = (128, 128))
    tmp_img = image.img_to_array(tmp_img)
    tmp_img = np.expand_dims(tmp_img, axis = 0)


    tmp_img /=255.
    plt.imshow(tmp_img[0])
    plt.show()
    # predict
    result = model3.predict(tmp_img)
    train_generator.class_indices
    print('Actual value: ',i,'\tPredicted value: ', np.argmax(result))
    i+=1

print('Total number of images for "testing":')
test_generator = test_datagen.flow_from_directory(test_dir,
                                                    target_size = (128, 128),
                                                    batch_size = 32,
                                                    class_mode = "categorical",
                                                    shuffle=False)


```




Actual value: 0 Predicted value: 0




Actual value: 1 Predicted value: 0



Actual value: 2 Predicted value: 2



Actual value: 3 Predicted value: 3



Actual value: 4 Predicted value: 4
Total number of images for "testing":
Found 110 images belonging to 5 classes.

```
In [165]: # MODEL 4
# -----
# ReLU activation function
# Rmsprop optimizer
# 4 Conv
# 1 Batch Norm
# 4 MaxPool
# 1 Dropout
# 3 Dense
# 1 Flatten
# Batchsize 32

model4 = tensorflow.keras.Sequential()
model4.add(layers.Conv2D(32, (3,3), activation='relu',
                        input_shape = (128, 128, 3),
                        kernel_initializer = 'glorot_normal',
                        bias_initializer = 'zeros'))
model4.add(layers.BatchNormalization())
model4.add(layers.MaxPooling2D((2,2)))
model4.add(layers.Conv2D(64, (3,3), activation='relu'))
model4.add(layers.MaxPooling2D((2,2)))
model4.add(layers.Conv2D(128, (3,3), activation='relu'))
model4.add(layers.MaxPooling2D((2,2)))
model4.add(layers.Conv2D(128, (3,3), activation='relu'))
model4.add(layers.MaxPooling2D((2,2)))
model4.add(layers.Flatten())
model4.add(layers.Dropout(.5))
model4.add(layers.Dense(128, activation='relu',
                        kernel_initializer = 'glorot_normal',
                        bias_initializer = 'zeros'))
model4.add(layers.Dense(512, activation='relu'))
model4.add(layers.Dense(5, activation='softmax'))
model4.summary()
```

Model: "sequential_32"

| Layer (type) | Output Shape | Param # |
|--|----------------------|---------|
| ===== | | |
| conv2d_127 (Conv2D) | (None, 126, 126, 32) | 896 |
| batch_normalization_32 (Batch Normalization) | (None, 126, 126, 32) | 128 |
| max_pooling2d_127 (MaxPooling2D) | (None, 63, 63, 32) | 0 |
| conv2d_128 (Conv2D) | (None, 61, 61, 64) | 18496 |
| max_pooling2d_128 (MaxPooling2D) | (None, 30, 30, 64) | 0 |
| conv2d_129 (Conv2D) | (None, 28, 28, 128) | 73856 |
| max_pooling2d_129 (MaxPooling2D) | (None, 14, 14, 128) | 0 |
| conv2d_130 (Conv2D) | (None, 12, 12, 128) | 147584 |
| max_pooling2d_130 (MaxPooling2D) | (None, 6, 6, 128) | 0 |
| flatten_32 (Flatten) | (None, 4608) | 0 |
| dropout_32 (Dropout) | (None, 4608) | 0 |

| | | |
|---------------------------|-------------|--------|
| dense_94 (Dense) | (None, 128) | 589952 |
| dense_95 (Dense) | (None, 512) | 66048 |
| dense_96 (Dense) | (None, 5) | 2565 |
| ===== | | |
| Total params: 899,525 | | |
| Trainable params: 899,461 | | |
| Non-trainable params: 64 | | |

```
In [166]: # SGD optimizer with Nesterov momentum
sgd = optimizers.SGD(lr=0.001, decay=1e-2, momentum=0.9, nesterov=True)

model4.compile(loss = "categorical_crossentropy",
               optimizer=sgd,
               metrics=["acc"])
```

In [167]: *# Training model with adagrad, 50 epochs and shuffling data*

```
batch_size=32
```

```
history_sgd = model4.fit_generator(train_generator,  
                                   steps_per_epoch=int(444/batch_size),  
                                   epochs=30,  
                                   validation_data=validation_generator,  
                                   validation_steps=int(110/batch_size),  
                                   shuffle=True)
```

WARNING:tensorflow:sample_weight modes were coerced from

```
...  
to  
['...']
```

WARNING:tensorflow:sample_weight modes were coerced from

```
...  
to  
['...']
```

Train for 13 steps, validate for 3 steps

Epoch 1/30

13/13 [=====] - 8s 639ms/step - loss: 1.7144 - acc: 0.1432 -
val_loss: 1.6109 - val_acc: 0.2396

Epoch 2/30

13/13 [=====] - 8s 596ms/step - loss: 1.6335 - acc: 0.2039 -
val_loss: 1.6086 - val_acc: 0.2500

Epoch 3/30

13/13 [=====] - 8s 595ms/step - loss: 1.6168 - acc: 0.1893 -
val_loss: 1.6049 - val_acc: 0.2292

Epoch 4/30

13/13 [=====] - 8s 612ms/step - loss: 1.6093 - acc: 0.2160 -
val_loss: 1.6054 - val_acc: 0.3125

Epoch 5/30

13/13 [=====] - 8s 594ms/step - loss: 1.6158 - acc: 0.2184 -
val_loss: 1.6046 - val_acc: 0.3021

Epoch 6/30

13/13 [=====] - 8s 596ms/step - loss: 1.6140 - acc: 0.2184 -
val_loss: 1.6026 - val_acc: 0.2917

Epoch 7/30

13/13 [=====] - 8s 598ms/step - loss: 1.6045 - acc: 0.2112 -
val_loss: 1.6031 - val_acc: 0.3021

Epoch 8/30

13/13 [=====] - 8s 599ms/step - loss: 1.6070 - acc: 0.2330 -
val_loss: 1.6041 - val_acc: 0.2396

Epoch 9/30

13/13 [=====] - 8s 605ms/step - loss: 1.6045 - acc: 0.2428 -
val_loss: 1.6018 - val_acc: 0.2812

Epoch 10/30

13/13 [=====] - 8s 595ms/step - loss: 1.6037 - acc: 0.2330 -
val_loss: 1.6007 - val_acc: 0.3021

Epoch 11/30

13/13 [=====] - 8s 593ms/step - loss: 1.6040 - acc: 0.2306 -
val_loss: 1.6016 - val_acc: 0.2604

Epoch 12/30

13/13 [=====] - 8s 627ms/step - loss: 1.6137 - acc: 0.2160 -
val_loss: 1.6004 - val_acc: 0.3021

Epoch 13/30

13/13 [=====] - 8s 623ms/step - loss: 1.5989 - acc: 0.2257 -
val_loss: 1.5974 - val_acc: 0.3333

Epoch 14/30


```

13/13 [=====] - 8s 602ms/step - loss: 1.5897 - acc: 0.2621 -
val_loss: 1.5950 - val_acc: 0.3021
Epoch 15/30
13/13 [=====] - 8s 595ms/step - loss: 1.5809 - acc: 0.2961 -
val_loss: 1.5959 - val_acc: 0.2604
Epoch 16/30
13/13 [=====] - 8s 609ms/step - loss: 1.5949 - acc: 0.2330 -
val_loss: 1.5924 - val_acc: 0.3333
Epoch 17/30
13/13 [=====] - 8s 596ms/step - loss: 1.5795 - acc: 0.2816 -
val_loss: 1.5909 - val_acc: 0.3333
Epoch 18/30
13/13 [=====] - 8s 600ms/step - loss: 1.5852 - acc: 0.2718 -
val_loss: 1.5863 - val_acc: 0.3750
Epoch 19/30
13/13 [=====] - 8s 595ms/step - loss: 1.5711 - acc: 0.2694 -
val_loss: 1.5847 - val_acc: 0.3750
Epoch 20/30
13/13 [=====] - 8s 627ms/step - loss: 1.5911 - acc: 0.2379 -
val_loss: 1.5827 - val_acc: 0.3333
Epoch 21/30
13/13 [=====] - 8s 606ms/step - loss: 1.5737 - acc: 0.2694 -
val_loss: 1.5809 - val_acc: 0.3438
Epoch 22/30
13/13 [=====] - 8s 596ms/step - loss: 1.5683 - acc: 0.2816 -
val_loss: 1.5740 - val_acc: 0.3646
Epoch 23/30
13/13 [=====] - 8s 594ms/step - loss: 1.5559 - acc: 0.3058 -
val_loss: 1.5666 - val_acc: 0.3958
Epoch 24/30
13/13 [=====] - 8s 596ms/step - loss: 1.6032 - acc: 0.2282 -
val_loss: 1.5697 - val_acc: 0.3438
Epoch 25/30
13/13 [=====] - 8s 602ms/step - loss: 1.5769 - acc: 0.2816 -
val_loss: 1.5649 - val_acc: 0.3854
Epoch 26/30
13/13 [=====] - 8s 599ms/step - loss: 1.5547 - acc: 0.3083 -
val_loss: 1.5579 - val_acc: 0.4062
Epoch 27/30
13/13 [=====] - 8s 605ms/step - loss: 1.5698 - acc: 0.2646 -
val_loss: 1.5505 - val_acc: 0.4062
Epoch 28/30
13/13 [=====] - 8s 620ms/step - loss: 1.5599 - acc: 0.3058 -
val_loss: 1.5503 - val_acc: 0.3750
Epoch 29/30
13/13 [=====] - 9s 681ms/step - loss: 1.5570 - acc: 0.2743 -
val_loss: 1.5543 - val_acc: 0.3229
Epoch 30/30
13/13 [=====] - 9s 671ms/step - loss: 1.5566 - acc: 0.2694 -
val_loss: 1.5486 - val_acc: 0.3542

```

```

In [51]: # -----
# SAVE THE MODEL |
# -----
model4.save('DIP_Proj_modelRSM2.h5')

```

```

In [52]: # -----
# Visualizing Train/Validation Loss & Accuracy /
# -----

acc_sgd = history_sgd.history['acc']
val_acc_sgd = history_sgd.history['val_acc']
loss_sgd = history_sgd.history['loss']
val_loss_sgd = history_sgd.history['val_loss']

epochs_sgd = range(1, len(acc_sgd) + 1)


# Plot of accuracy
plt.plot(epochs_sgd, acc_sgd, color='blue', label='Training Acc')
plt.plot(epochs_sgd, val_acc_sgd, color='red', label='Validation Acc')
plt.title('Training and Validation Accuracy (RSM)')
plt.legend()


plt.figure()

# Plot of loss
plt.plot(epochs_sgd, loss_sgd, color='blue', label='Training Loss')
plt.plot(epochs_sgd, val_loss_sgd, color='red', label='Validation Loss')
plt.title('Training and Validation Loss (RSM)')
plt.legend()
plt.figure()

plt.show()

```





<Figure size 432x288 with 0 Axes>

```

In [53]: # -----
# ROC/AUC Score |
# -----

from sklearn.preprocessing import LabelBinarizer

# set plot figure size
fig, c_ax = plt.subplots(1,1, figsize = (12, 8))

def multiclass_roc_auc_score(y_test, y_pred, average="macro"):
    lb = LabelBinarizer()
    lb.fit(y_test)
    y_test = lb.transform(y_test)
    y_pred = lb.transform(y_pred)


    for (idx, c_label) in enumerate(target_names): # target_names: no of the labels
        fpr, tpr, thresholds = roc_curve(y_test[:,idx].astype(int), y_pred[:,idx])
        c_ax.plot(fpr, tpr, label = '%s (AUC:%0.2f)' % (c_label, auc(fpr, tpr)))
    c_ax.plot(fpr, fpr, 'b-', label = 'Random Guessing')
    return roc_auc_score(y_test, y_pred, average=average)

validation_generator.reset() # resetting generator
y_pred = model3.predict_generator(validation_generator, verbose = True)
y_pred = np.argmax(y_pred, axis=1)
multiclass_roc_auc_score(validation_generator.classes, y_pred)

```

4/4 [=====] - 0s 111ms/step

Out[53]: 0.5



```
In [54]: # -----
# Classification Report /
# -----

print('Classification Report for Model with RSM')
target_names = ['1_Neutral', '2_Smiling', '3_Sleepy', '4_Surprise', '5_Sunglasses']
print(classification_report(validation_generator.classes, y_pred, target_names=target_names))
```

```
Classification Report for Model with RSM
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 1_Neutral | 0.00 | 0.00 | 0.00 | 22 |
| 2_Smiling | 0.00 | 0.00 | 0.00 | 22 |
| 3_Sleepy | 0.00 | 0.00 | 0.00 | 22 |
| 4_Surprise | 0.21 | 1.00 | 0.34 | 22 |
| 5_Sunglasses | 0.00 | 0.00 | 0.00 | 22 |
| accuracy | | | 0.20 | 110 |
| macro avg | 0.04 | 0.20 | 0.07 | 110 |
| weighted avg | 0.04 | 0.20 | 0.07 | 110 |

C:\Users\User\anaconda3\envs\tensorflow\lib\site-packages\sklearn\metrics_classification.py:1221: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
In [55]: # -----
# Confusion Matrix /
# -----

num_of_train_samples = 444
#num_of_test_samples = 416 # steps per epoch
batch_size=32
steps_per_epoch=num_of_train_samples // batch_size
#validation_generator.reset()

Y_pred = model3.predict_generator(validation_generator, steps_per_epoch)
y_pred = np.argmax(Y_pred, axis=1)
print('Confusion Matrix for Model with RSM')
print(confusion_matrix(validation_generator.classes, y_pred))
```

WARNING:tensorflow:Your input ran out of data; interrupting training. Make sure that your dataset or generator can generate at least `steps_per_epoch * epochs` batches (in this case, 13 batches). You may need to use the repeat() function when building your dataset.

Confusion Matrix for Model with RSM

```
[[ 0  0  0 20  2]
 [ 0  0  0 21  1]
 [ 0  0  0 21  1]
 [ 0  0  0 22  0]
 [ 0  0  0 22  0]]
```

```

In [56]: # -----
# Displaying 12 Images with the Prediction |
# -----

# 1 image from each class

labels = ['1_Neutral', '2_Smiling', '3_Sleepy', '4_Surprise', '5_Sunglasses']


i=0
for names in labels:
    pathToFolder = test_dir + '/' + names + '/'
    fnames = [os.path.join(pathToFolder, fname) for
               fname in os.listdir(pathToFolder)]

    # generate random number btwn (0,30)
    randNum = random.randint(0, 20)
    img_path = fnames[randNum]
    tmp_img = image.load_img(img_path, target_size = (128, 128))
    tmp_img = image.img_to_array(tmp_img)
    tmp_img = np.expand_dims(tmp_img, axis = 0)

    tmp_img /=255.
    plt.imshow(tmp_img[0])
    plt.show()
    # predict
    result = model3.predict(tmp_img)
    train_generator.class_indices
    print('Actual value: ',i,'\tPredicted value: ', np.argmax(result))
    i+=1


print('Total number of images for "testing":')
test_generator = test_datagen.flow_from_directory(test_dir,
                                                    target_size = (128, 128),
                                                    batch_size = 32,
                                                    class_mode = "categorical",
                                                    shuffle=False)

```




Actual value: 0


Predicted value: 3




Actual value: 1 Predicted value: 3



Actual value: 2 Predicted value: 3



Actual value: 3 Predicted value: 3



Actual value: 4 Predicted value: 3
Total number of images for "testing":
Found 110 images belonging to 5 classes.

In []: